

# **Vector Databases & Hybrid Search**

## **Semantic + Keyword Search for Business Intelligence**

**FAISS, Embeddings, and Real-World Applications**

# Module Overview

## What You'll Learn

- **Hybrid search concepts** - Combining semantic and keyword search
- **Embeddings** - Converting text to numbers for AI understanding
- **FAISS** - Vector database for similarity search
- **FTS5** - SQLite's full-text search capabilities
- **Weight modification** - Tuning search results for your use case

# What is Hybrid Search?

## The Best of Both Worlds

Hybrid search combines **two different search approaches**:

1. **Keyword Search (FTS5 + BM25)** - Finds exact matches and phrases
2. **Semantic Search (FAISS + Embeddings)** - Understands meaning and context

**Why Hybrid?** Sometimes you want exact matches ("social media strategy"), sometimes you want related concepts ("marketing approach" → finds "brand strategy").

# Understanding the Components

## 1. Embeddings: Converting Text to Numbers

### What are embeddings?

- Embeddings convert text into a list of numbers (vectors)
- Similar text gets similar numbers
- Our app uses OpenAI's `text-embedding-3-small` model (1536 dimensions)

# Embeddings Example

## How It Works

```
# Text: "marketing strategy"  
# Embedding: [0.1, -0.3, 0.8, 0.2, ...] (1536 numbers)  
  
# Text: "brand approach"  
# Embedding: [0.2, -0.2, 0.7, 0.3, ...] (similar numbers = similar meaning)
```

**Similar numbers = similar meaning**

# Embeddings Implementation

## OpenAI Integration

```
# services/embedding_service.py
async def generate_embedding(self, text: str) -> List[float]:
    response = await self.client.embeddings.create(
        model="text-embedding-3-small",
        input=text
    )
    return response.data[0].embedding # Returns 1536 numbers
```

**1536-dimensional vectors for semantic understanding**

# Understanding the Components

## 2. FAISS: Vector Database

### What is FAISS?

- Facebook AI Similarity Search
- Stores and searches through vectors efficiently
- Finds similar vectors quickly (cosine similarity)

# Vector Database Comparison

## Choosing the Right Solution

Feature	FAISS	ChromaDB	Pinecone	PostgreSQL + pgvector
Setup	Simple, local	Easy, local/cloud	Cloud service	Moderate, local/cloud
Cost	Free	Free/paid	Paid	Free/paid
Performance	Very fast	Good	Good	Good
Scalability	Self-managed	Self-managed	Managed	Self-managed



# FAISS Implementation

## Current Setup

```
# services/hybrid_search_service.py
class HybridSearchService:
    def __init__(self):
        self.faiss_index = faiss.IndexFlatIP(1536) # Inner product index
        self.embedding_dimension = 1536
```

Inner product index for cosine similarity

# Switching Vector Databases

## Easy Migration Strategy

The vector database implementation is abstracted in our `HybridSearchService` class:

1. **Create a new vector database adapter** (e.g., `PostgreSQLVectorAdapter`)
2. **Implement the same interface** as the current FAISS implementation
3. **Update the service initialization** to use the new adapter
4. **Migrate existing vector data** to the new database

# PostgreSQL + pgvector Example

## Alternative Implementation

```
# New adapter class
class PostgreSQLVectorAdapter:
    def __init__(self, connection_string: str):
        self.conn = psycopg2.connect(connection_string)

    def add_vectors(self, vectors: List[List[float]], ids: List[int]):
        # Insert vectors into PostgreSQL with pgvector
        pass

    def search(self, query_vector: List[float], k: int = 10):
        # Use pgvector's similarity search
        pass

# Update HybridSearchService
class HybridSearchService:
    def __init__(self, vector_db_type: str = "faiss"):
```

# Understanding the Components

## 3. FTS5: Full-Text Search

### What is FTS5?

- SQLite's Full-Text Search extension
- Searches through text content efficiently
- Uses BM25 ranking algorithm

# FTS5 Implementation

## Virtual Table Setup

```
-- Virtual table for searchable content
CREATE VIRTUAL TABLE chunks_fts USING fts5(
    content,
    chunk_type,
    content='chunks',
    content_rowid='rowid'
);
```

## BM25 ranking for keyword relevance

# Understanding Weight Modification

## The Hybrid Score Formula

```
hybrid_score = (bm25_weight × BM25_score) + (cosine_weight × cosine_similarity_score)
```

### Default Weights:

- `bm25_weight = 0.35` (35% keyword search)
- `cosine_weight = 0.65` (65% semantic search)

# Weight Modification Scenarios

## Scenario 1: More Keyword Focus

```
# For exact term matching (legal documents, code)  
bm25_weight = 0.7    # 70% keyword  
cosine_weight = 0.3  # 30% semantic
```

**Use case: Legal documents, code repositories**

# Weight Modification Scenarios

## Scenario 2: More Semantic Focus

```
# For creative content (marketing, research)  
bm25_weight = 0.2    # 20% keyword  
cosine_weight = 0.8  # 80% semantic
```

**Use case: Marketing content, research papers**



# Weight Modification Scenarios

## Scenario 3: Balanced Search

```
# For general content  
bm25_weight = 0.5    # 50% keyword  
cosine_weight = 0.5  # 50% semantic
```

**Use case: General business documents**

# Testing Different Weights

## API Integration

```
# API call with custom weights
```

```
GET /search/hybrid?q=marketing&bm25_weight=0.2&cosine_weight=0.8
```

**Easy experimentation with different weight combinations**

# Learning Path

## Phase 1: Understanding the Basics

1. **Study the embedding service** ( `services/embedding_service.py` )
2. **Examine chunk creation** ( `services/chunking_service.py` )
3. **Test different search methods** via the UI at `/search`

# Learning Path

## Phase 2: Experimenting with Weights

1. **Modify weights** in the search UI
2. **Compare results** between keyword, semantic, and hybrid
3. **Test with different content types**

# Learning Path

## Phase 3: Extending the System

1. **Add new content types** (PDFs, images)
2. **Implement custom scoring** algorithms
3. **Add search filters** and faceted search

# Learning Path

## Phase 4: Advanced Features

1. **Add search suggestions** and autocomplete
2. **Implement search analytics** and user behavior tracking
3. **Add real-time search** updates

# Architecture Deep Dive

## Data Flow

1. User types search query
2. Query gets embedded (text → numbers)
3. FAISS finds similar vectors
4. FTS5 finds keyword matches
5. Results get combined with weights
6. Final ranked results returned

# Key Files to Study

## Implementation Details

- `services/hybrid_search_service.py` - Core hybrid search logic
- `services/embedding_service.py` - OpenAI integration
- `services/chunking_service.py` - Content processing
- `routes/marketing.py` - Search API endpoints
- `templates/search.html` - Frontend interface



# Extending for Real-World Use

## Adding New Content Types

```
# Example: Add PDF support  
class PDFChunkingService:  
    def chunk_pdf(self, pdf_path: str) -> List[Chunk]:  
        # Extract text from PDF  
        # Split into chunks  
        # Generate embeddings  
        pass
```

**Extensible architecture for any content type**

# Custom Scoring Algorithms

## Example: Boost Recent Content

```
# Example: Boost recent content  
def custom_score(self, chunk, base_score):  
    recency_boost = self.calculate_recency_boost(chunk.created_at)  
    return base_score * recency_boost
```

Customize ranking for your business needs

# Performance Optimization

## Caching Frequent Searches

```
# Example: Caching frequent searches  
@lru_cache(maxsize=1000)  
def cached_embedding(self, text: str):  
    return self.generate_embedding(text)
```

Improve performance with intelligent caching

# Monitoring and Analytics

## Search Metrics to Track

- **Query performance** (response time)
- **Result relevance** (click-through rates)
- **Popular queries** and search patterns
- **Search method effectiveness** (hybrid vs keyword vs semantic)

# Analytics Implementation

## Tracking Search Metrics

```
# Add to search endpoints  
async def track_search(query: str, method: str, results_count: int):  
    # Log search metrics  
    # Update analytics database  
    pass
```

## Data-driven search optimization

# Troubleshooting Common Issues

## Issue 1: No Search Results

- **Check:** Are chunks being created?
- **Check:** Are embeddings being generated?
- **Check:** Is FAISS index built?

# Troubleshooting Common Issues

## Issue 2: Poor Search Quality

- **Try:** Adjusting weights
- **Try:** Different chunk sizes
- **Try:** Different embedding models

# Troubleshooting Common Issues

## Issue 3: Slow Performance

- **Check:** FAISS index size
- **Check:** Database query performance
- **Consider:** Adding caching



# Corporate Database Use Case

## 10,000 Customer Analysis

This architecture is perfect for analyzing a corporate database of 10,000 customers/stakeholders for cross-selling opportunities.

# Why This Architecture Works for Customer Analysis

## Key Benefits

1. **Semantic Understanding** - Find customers with similar interests even if they used different words
2. **Hybrid Search** - Combine exact matches (company names, industries) with semantic matches (similar business needs)
3. **Scalable** - FAISS handles 10,000+ records efficiently
4. **Real-time** - Update customer profiles and immediately searchable

# Customer Data Structure

## Implementation for Customer Analysis

```
class CustomerProfile:  
    company_name: str  
    industry: str  
    business_needs: str # Rich text description  
    current_products: List[str]  
    pain_points: str  
    company_size: str  
    location: str  
    engagement_level: str # webinar, whitepaper, trial, etc.
```

# Search Use Cases

## Customer Analysis Queries

- **"Find companies like Acme Corp"** → Semantic search finds similar business profiles
- **"Manufacturing companies needing automation"** → Hybrid search finds exact + related matches
- **"SaaS companies with 50-200 employees"** → Filtered semantic search

# Cross-Selling Opportunities

## Sales Team Queries

*# Example search queries for sales teams:*

```
queries = [  
    "companies using our CRM who might need marketing automation",  
    "enterprise clients who haven't tried our analytics product",  
    "SMB customers ready for enterprise features",  
    "companies in healthcare needing compliance tools"  
]
```

# Enhanced Features for Sales

## Advanced Customer Analysis

- **Customer Similarity** - "Show me 10 companies most similar to [current customer]"
- **Gap Analysis** - "What products do similar companies use that this customer doesn't?"
- **Engagement Scoring** - Boost customers with high engagement in search results
- **Timing Analysis** - Find customers who engaged 6+ months ago but haven't converted

# Performance at Scale

## Handling Large Datasets

- **10,000 customers** - FAISS handles this easily (can scale to millions)
- **Real-time updates** - New customer data immediately searchable
- **Fast queries** - Sub-second response times even with complex searches

# Business Value

## ROI for Customer Analysis

- **Increase Revenue** - Find cross-selling opportunities automatically
- **Improve Targeting** - Better customer segmentation and personalization
- **Save Time** - Sales teams find relevant prospects instantly
- **Data-Driven** - Use actual customer behavior and needs, not just demographics



# Further Learning Resources

## Additional Study Materials

- **FAISS Documentation:** <https://faiss.ai/>
- **OpenAI Embeddings Guide:** <https://platform.openai.com/docs/guides/embeddings>
- **SQLite FTS5:** <https://www.sqlite.org/fts5.html>
- **Vector Search Best Practices:** <https://weaviate.io/blog/vector-search-best-practices>

# Next Steps

## Immediate Actions

1. **Experiment** with the search interface
2. **Modify weights** and see how results change
3. **Add your own content** and test search quality
4. **Extend the system** with new features
5. **Deploy** to a real environment

# Key Takeaways

## What Makes Hybrid Search Powerful

1. **Combines best of both worlds** - Exact matches + semantic understanding
2. **Highly tunable** - Adjust weights for your specific use case
3. **Scalable** - Handles large datasets efficiently
4. **Real-world ready** - Perfect for business applications
5. **Extensible** - Easy to add new content types and features

# Key Takeaways

## Business Applications

- **Customer analysis** - Find similar customers and cross-selling opportunities
- **Document search** - Intelligent content discovery
- **Knowledge management** - Organize and find information efficiently
- **Research** - Discover related concepts and ideas
- **Content recommendation** - Suggest relevant content to users

# Ready to Build?

## Start Learning

**The best way to learn is by doing!**

Start with small modifications and gradually build more complex features.

**Let's create something amazing! 🚀**